



IBM Developer
SKILLS NETWORK

Winning Space Race with Data Science

Rajesh Karunakaran
March 9, 2022



Outline

- Executive Summary
- Introduction
- Methodology
- Results
- Conclusion
- Appendix

Executive Summary

- **Summary of methodologies**
 - Data Collection using API
 - Data Collection with Web Scraping
 - Data Wrangling
 - Exploratory Data Analysis with SQL
 - Exploratory Data Analysis using Data Visualization
 - Interactive Visual Analytics with Folium
 - Interactive Dashboard with Plotly Dash
 - Machine Learning Classification
- **Summary of all results**
 - Exploratory Data Analysis result
 - Interactive analytics in screenshots
 - Predictive Analysis result

Introduction

- **Project background and context**

There are several companies that are making space travel affordable for everyone. Perhaps the most successful of them is SpaceX, and one of the reasons is that their rocket launch is relatively inexpensive.

SpaceX advertises Falcon 9 rocket launches on its website, with a cost of 62 million dollars; other providers cost upward of 165 million dollars each, much of the savings is because SpaceX can reuse the first stage.

Therefore, we will predict if the Falcon 9 first stage will land successfully. If we can determine if the first stage will land, we can determine the cost of a launch. This information can be used if an alternate company wants to bid against SpaceX for a rocket launch.

- **Problems you want to find answers**

- Correlations between each rocket variables and successful landing rate
- Conditions to get the best results and ensure the best successful landing rate

Section 1

Methodology

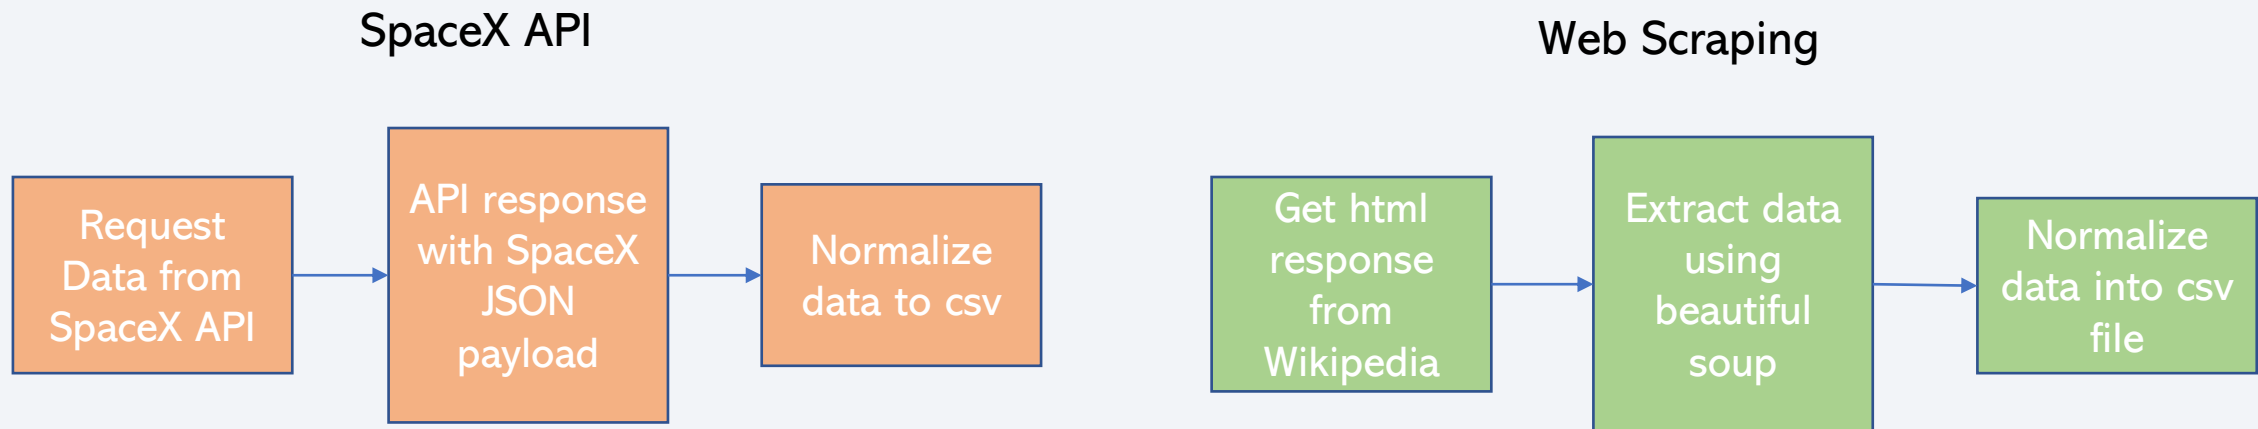
Methodology

Executive Summary

- Data collection methodology:
 - Data was collected using SpaceX API and web scraping from Wikipedia.
- Perform data wrangling
 - Convert outcomes into labels for training the supervised ML models.
- Perform exploratory data analysis (EDA) using visualization and SQL
- Perform interactive visual analytics using Folium and Plotly Dash
- Perform predictive analysis using classification models
 - Find best Hyperparameter for SVM, Decision Trees, KNN and Logistic Regression

Data Collection

- The data collection process includes 2 methods using API requests from the SpaceX API and web scraping data from a table in the Wikipedia.
 - SpaceX API Data Columns: FlightNumber, Date, BoosterVersion, PayloadMass, Orbit, LaunchSite, Outcome, Flights, GridFins, Reused, Legs, LandingPad, Block, ReusedCount, Serial, Longitude, Latitude ○
 - Wikipedia Web Scrape Data Columns: Flight No., Launch site, Payload, PayloadMass, Orbit, Customer, Launch outcome, Version Booster, Booster landing, Date, Time



Data Collection – SpaceX API

1. Requesting rocket launch data from SpaceX API

```
spacex_url="https://api.spacexdata.com/v4/launches/past"

response = requests.get(spacex_url)
```

2. Converting Response to a JSON file

```
data = pd.json_normalize(response.json())
```

3. Using custom functions to clean data

```
# Call getBoosterVersion
getBoosterVersion(data)
```

```
# Call getPayloadData
getPayloadData(data)
```

```
# Call getLaunchSite
getLaunchSite(data)
```

```
# Call getCoreData
getCoreData(data)
```

<https://github.com/rajesh16rk/IBM-SpaceX-Applied-Capstone-Project/blob/main/1.%20SpaceX%20Data%20Collection%20API.ipynb>

4. Combining the columns into a dictionary to create data frame

```
launch_dict = {'FlightNumber': list(data['flight_number']),
               'Date': list(data['date']),
               'BoosterVersion': BoosterVersion,
               'PayloadMass': PayloadMass,
               'Orbit': Orbit,
               'LaunchSite': LaunchSite,
               'Outcome': Outcome,
               'Flights': Flights,
               'GridFins': GridFins,
               'Reused': Reused,
               'Legs': Legs,
               'LandingPad': LandingPad,
               'Block': Block,
               'ReusedCount': ReusedCount,
               'Serial': Serial,
               'Longitude': Longitude,
               'Latitude': Latitude}
```

```
launch_df = pd.DataFrame.from_dict(launch_dict)
```

5. Filtering dataframe and exporting to a CSV

```
data_falcon9 = launch_df[launch_df['BoosterVersion'] == 'Falcon 9']
```

```
data_falcon9.to_csv('dataset_part_1.csv', index=False)
```


Data Collection - Scraping

1. Getting response from HTML

```
html_data = requests.get(static_url).text
```

2. Creating a BeautifulSoup object

```
soup = BeautifulSoup(html_data, 'html5lib')
```

3. Finding all tables and assigning the result to a list

```
html_tables = soup.find_all('table')
```

4. Extracting column name one by one

```
column_names = []

for row in first_launch_table.find_all('th'):
    name = extract_column_from_header(row)
    if (name != None and len(name) > 0):
        column_names.append(name)
```

<https://github.com/rajesh16rk/IBM-SpaceX-Applied-Capstone-Project/blob/main/2.%20SpaceX%20Data%20Collection%20Web%20Scraping.ipynb>

5. Creating an empty dictionary with keys

```
launch_dict = dict.fromkeys(column_names)

# Remove an irrelevant column
del launch_dict['Date and time ( )']

launch_dict['Flight No.'] = []
launch_dict['Launch site'] = []
launch_dict['Payload'] = []
launch_dict['Payload mass'] = []
launch_dict['Orbit'] = []
launch_dict['Customer'] = []
launch_dict['Launch outcome'] = []

# Added some new columns
launch_dict['Version Booster'] = []
launch_dict['Booster landing'] = []
launch_dict['Date'] = []
launch_dict['Time'] = []
```

6. Filling up the launch_dict with launch records (Too long to put in here, so please refer to the notebook)

7. Creating a Dataframe and exporting it to a CSV

```
df = pd.DataFrame(launch_dict)

df.to_csv('spacex_web_scraped.csv', index=False)
```

Data Wrangling

1. Calculating the number of launches at each site

```
df['LaunchSite'].value_counts()
```

2. Calculating the number and occurrence of each orbit

```
df.Orbit.value_counts()
```

3. Calculating the number and occurrence of mission outcome per orbit type

```
landing_outcomes = df.Outcome.value_counts()
```

4. Creating a landing outcome label from Outcome column

```
landing_class = []
for outcome in df.Outcome:
    if outcome in bad_outcomes:
        landing_class.append(0)
    else:
        landing_class.append(1)

df['Class'] = landing_class
```

5. Calculating the success rate for every landing in dataset

```
df["Class"].mean()
0.6666666666666666
```

6. Exporting dataset to a CSV

```
df.to_csv("dataset_part_2.csv", index=False)
```

<https://github.com/rajesh16rk/IBM-SpaceX-Applied-Capstone-Project/blob/main/3.%20SpaceX%20Data%20Wrangling.ipynb>

EDA with Data Visualization

- Scatter chart:

- Flight Number vs. Launch Site
- Payload vs. Launch Site
- Flight Number vs. Orbit Type
- Payload vs. Orbit Type

<https://github.com/rajesh16rk/IBM-SpaceX-Applied-Capstone-Project/blob/main/5.%20SpaceX%20EDA%20with%20pandas%20and%20matplotlib.ipynb>

- A scatter plot shows how much one variable is affected by another. The relationship between two variables is called a correlation. This plot is generally composed of large data bodies.

- Bar chart:

- Orbit Type vs. Success Rate
- A Bar chart makes it easy to compare datasets between multiple groups at a glance. One axis represents a category and the other axis represents a discrete value. The purpose of this chart is to indicate the relationship between the two axes.

- Line chart:

- Year vs. Success Rate
- A Line chart shows data variables and trends very clearly and helps predict the results of data that has not yet been recorded.

EDA with SQL

- Loading the dataset into the corresponding table in a IBM DB2 database, and executing SQL queries to answer the following questions:
 - Displaying the names of the unique launch sites in the space mission
 - Displaying 5 records where launch sites begin with the string 'CCA'
 - Displaying the total payload mass carried by boosters launched by NASA (CRS)
 - Displaying average payload mass carried by booster version F9 v1.1
 - Listing the date when the first successful landing outcome in ground pad was achieved
 - Listing the names of the boosters which have success in drone ship and have payload mass greater than 4000 but less than 6000
 - Listing the total number of successful and failure mission outcomes
 - Listing the names of the booster_versions which have carried the maximum payload mass
 - Listing the failed landing_outcomes in drone ship, their booster versions, and launch site names for in year 2015
 - Ranking the count of landing outcomes (such as Failure (drone ship) or Success (ground pad)) between the date 2010-06-04 and 2017-03-20, in descending order

<https://github.com/rajesh16rk/IBM-SpaceX-Applied-Capstone-Project/blob/main/4.%20SpaceX%20EDA%20with%20SQL.ipynb>

Build an Interactive Map with Folium

- Objects created and added to a folium map:
 - Markers that show all launch sites on a map
 - Markers that show the success/failed launches for each site on the map
 - Lines that show the distances between a launch site to its proximities
- By adding these objects, following geographical patterns about launch sites are found:
 - Are launch sites in close proximity to railways? Yes
 - Are launch sites in close proximity to highways? Yes
 - Are launch sites in close proximity to coastline? Yes
 - Do launch sites keep certain distance away from cities? Yes

<https://github.com/rajesh16rk/IBM-SpaceX-Applied-Capstone-Project/blob/main/6.%20SpaceX%20Interactive%20Visual%20Analytics%20with%20Folium.ipynb>

Build a Dashboard with Plotly Dash

- The dashboard application contains a pie chart and a scatter point chart.
 - *Pie chart*
 - For showing total success launches by sites
 - This chart can be selected to indicate a successful landing distribution across all launch sites or to indicate the success rate of individual launch sites.
 - *Scatter chart*
 - For showing the relationship between Outcomes and Payload mass (Kg) by different boosters
 - Has 2 inputs: All sites/individual site & Payload mass on a slider between 0 and 10000 kg
 - This chart helps determine how success depends on the launch point, payload mass, and booster version categories.

<https://github.com/rajesh16rk/IBM-SpaceX-Applied-Capstone-Project/blob/main/7.%20SpaceX%20Interactive%20Dashboard%20with%20Plotly%20Dash.py>

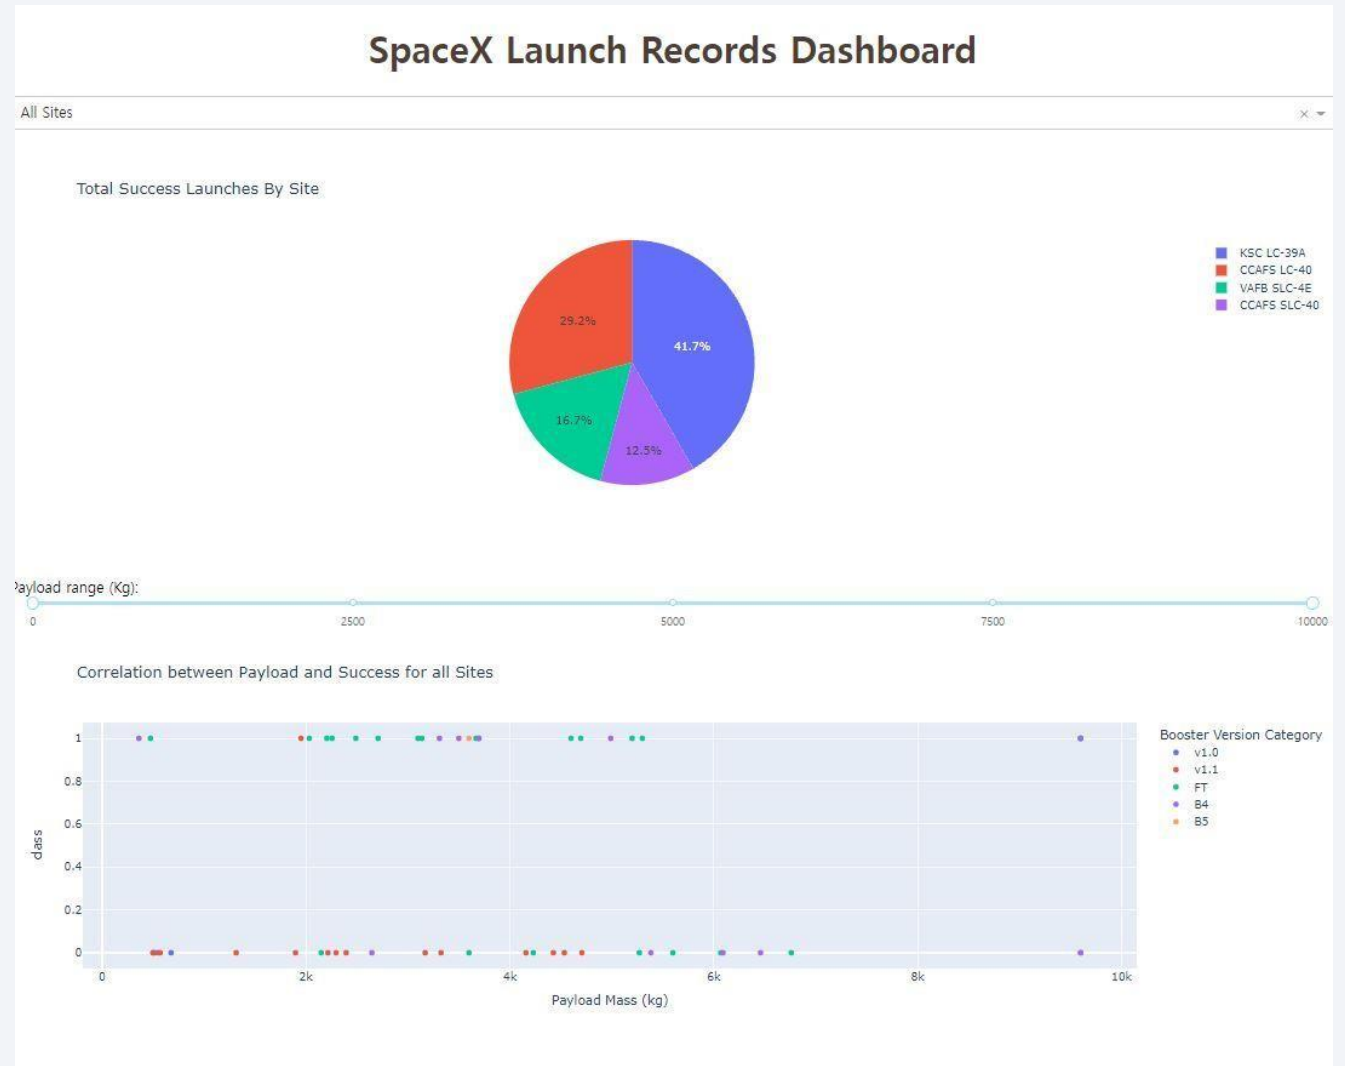
Predictive Analysis (Classification)

- Loaded the data using Numpy and pandas, transformed the data, split our data into training and testing
- Standardize the data and performed EDA
- Built different machine learning models and tune different hyperparameters using Grid Search CV
- Used accuracy as the metric for our model, improved the model using feature engineering and algorithm tuning
- Found the best performing classification model using the test data
- Found best Hyperparameter for SVM, Decision Trees, KNN and Logistic Regression

<https://github.com/rajesh16rk/IBM-SpaceX-Applied-Capstone-Project/blob/main/8.%20SpaceX%20MachineLearning%20Prediction.ipynb>

Results

- The right screenshot is a preview of the interactive Dashboard with Plotly Dash.
- The results of EDA with visualization, EDA with SQL, Interactive Map with Folium, and Interactive Dashboard will be shown in the next slides.
- Comparing the accuracy of the four models, all return the same accuracy of about 83.33% for test data




The background of the slide is an abstract composition. It features a dark blue base color. Overlaid on this are numerous diagonal streaks in shades of red and cyan. A faint, light blue grid pattern is also visible, particularly in the lower half of the image. The overall effect is dynamic and technological.

Section 2

Insights drawn from EDA

Flight Number vs. Launch Site

- Class 0 (blue) represents unsuccessful launch, and Class1 (orange) represents successful launch.
- This figure shows that the success rate increased as the number of flights increased.
- As the success rate has increased considerably since the 20th flight, this point seems to be a big breakthrough.
- Larger the flight amount at a launch site, the greater the success rate at a launch site




Payload vs. Launch Site

- Class 0 (blue) represents unsuccessful launch, and Class 1 (orange) represents successful launch
- At first glance, the larger payload mass, the higher the rocket's success rate

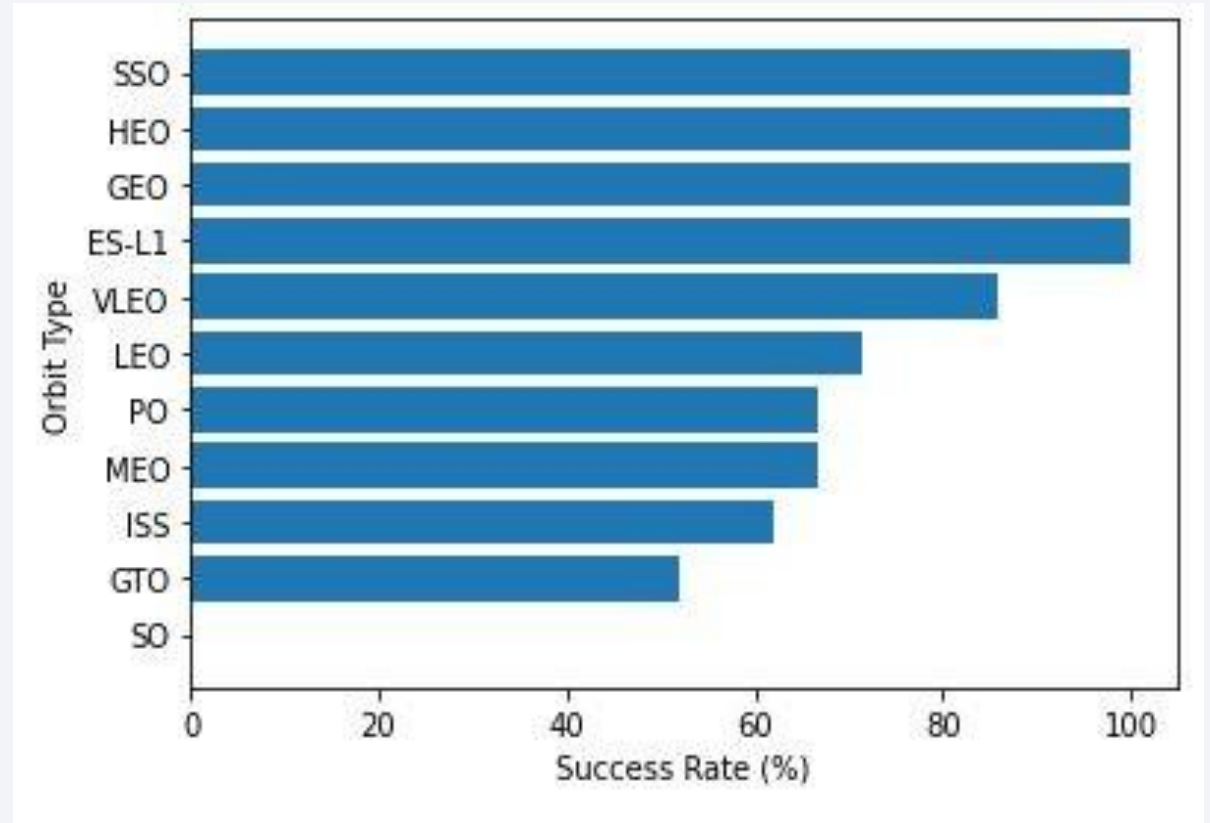


The greater the payload mass for launch site CCAFS SLC 40 the higher the success rate for the rocket.




Success Rate vs. Orbit Type

- Orbit types **SSO**, **HEO**, **GEO** and **ES-L1** have the highest success rates (100%).
- On the other hand, the success rate of orbit type **GTO** is only **50%**, and it is the lowest except for type **SO**, which recorded failure in a single attempt.




Flight Number vs. Orbit Type

- Class 0 (blue) represents unsuccessful launch, and Class 1 (orange) represents successful launch.
- In most cases, the launch outcome seems to be correlated with the flight number.
- On the other hand, in **GTO orbit**, there seems to be no relationship between flight numbers and success rate.
- SpaceX starts with LEO with a moderate success rate, and it seems that **VLEO**, which has a high success rate, is used the most in recent launches.



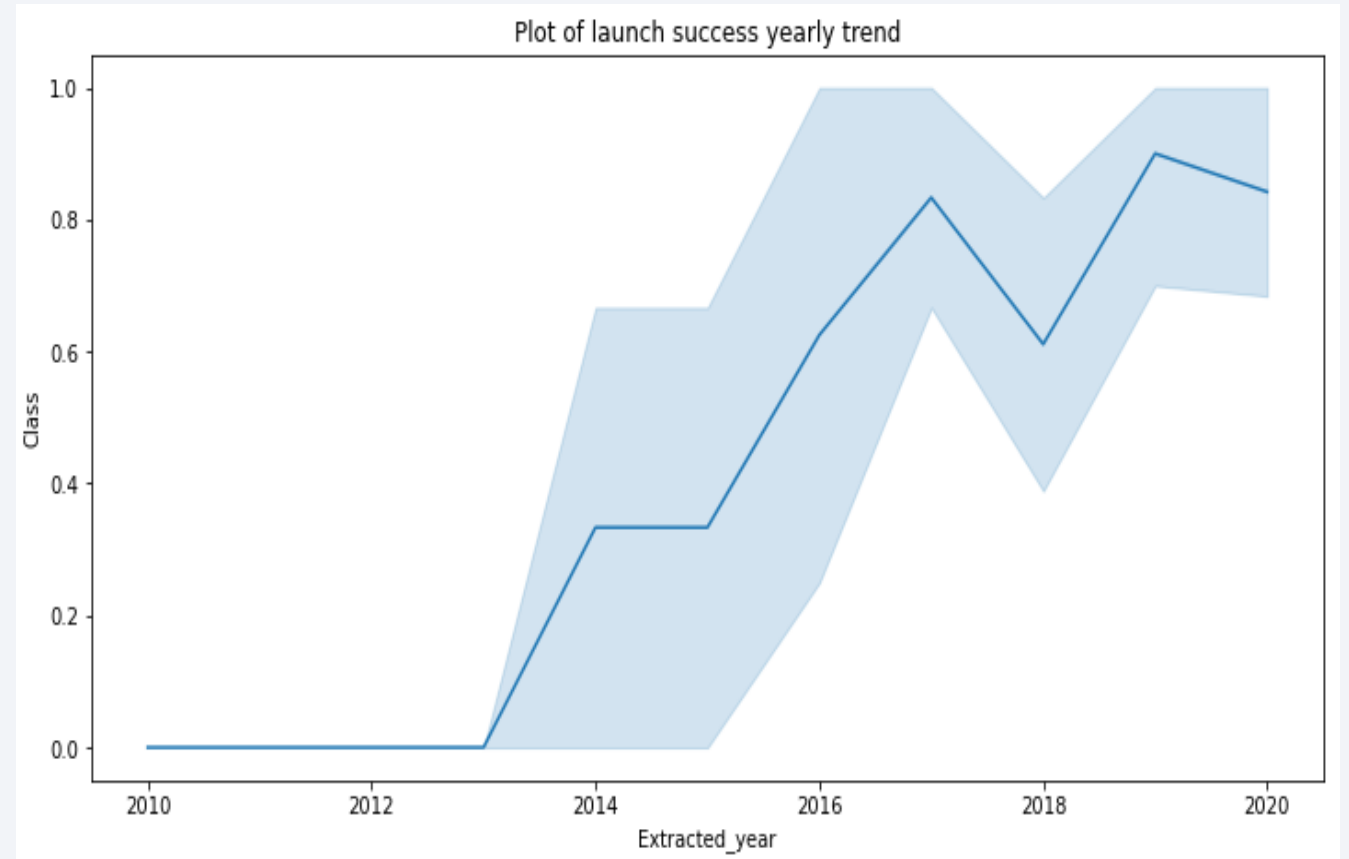
Payload vs. Orbit Type

- Class 0 (blue) represents unsuccessful launch, and Class 1 (orange) represents successful launch
- With heavy payloads the successful landing or positive landing rate are more for LEO and ISS orbits
- However, in the case of GTO, it is hard to distinguish between the positive landing rate and the negative landing because they are all gathered together



Launch Success Yearly Trend

- Since 2013, the success rate has continued to increase until 2020
- The rate decreased slightly in 2018
- Recently, it has shown a success rate of about 80%



All Launch Site Names

- SQL DISTINCT clause is used in the query, only unique values are displayed in the Launch Site column from the SpaceX table.
- There are four unique launch sites: CCAFS LC-40, CCAFS SLC-40, KSC LC-39A, VAFB SLC-4E

- Query

```
SELECT DISTINCT LAUNCH_SITE FROM  
SPACEXTBL
```

- Result

launch_site
CCAFS LC-40
CCAFS SLC-40
KSC LC-39A
VAFB SLC-4E

Launch Site Names Begin with 'CCA'

- Only five records of the SpaceX table were displayed using LIMIT 5 clause in the query.
- Using the LIKE operator and the percent sign (%) together, the Launch_Site name starting with CCA could be called.

- Query

```
SELECT * FROM SPACEXTBL
WHERE LAUNCH_SITE LIKE 'CCA%' LIMIT
5
```

- Result

DATE	time__utc_	booster_version	launch_site	payload	payload_mass__kg_	orbit	customer	mission_outcome	landing__outcome
2010-06-04	18:45:00	F9 v1.0 B0003	CCAFS LC-40	Dragon Spacecraft Qualification Unit	0	LEO	SpaceX	Success	Failure (parachute)
2010-12-08	15:43:00	F9 v1.0 B0004	CCAFS LC-40	Dragon demo flight C1, two CubeSats, barrel of Brouere cheese	0	LEO (ISS)	NASA (COTS) NRO	Success	Failure (parachute)
2012-05-22	07:44:00	F9 v1.0 B0005	CCAFS LC-40	Dragon demo flight C2	525	LEO (ISS)	NASA (COTS)	Success	No attempt
2012-10-08	00:35:00	F9 v1.0 B0006	CCAFS LC-40	SpaceX CRS-1	500	LEO (ISS)	NASA (CRS)	Success	No attempt
2013-03-01	15:10:00	F9 v1.0 B0007	CCAFS LC-40	SpaceX CRS-2	677	LEO (ISS)	NASA (CRS)	Success	No attempt

Total Payload Mass

- Calculated the total payload carried by boosters from NASA as 45596 using the query below
- Using the SUM() function to calculate the sum of column PAYLOAD_MASS_KG
- In the WHERE clause, filter the dataset to perform calculations only if Customer is NASA(CRS)

- Query

```
SELECT SUM(PAYLOAD_MASS_KG_)
        AS total_payload_mass_kg FROM
        SPACEXTBL
WHERE  CUSTOMER = 'NASA (CRS)'
```

- Result

total_payload_mass_kg
45596

Average Payload Mass by F9 v1.1

- Calculated the average payload mass carried by booster version F9 v1.1 as 2928.4
- Using the AVG() function to calculate the average value of column PAYLOAD_MASS_KG_
- In the WHERE clause, filter the dataset to perform calculations only if Booster_version is F9 v1.1

- Query

```
SELECT AVG(PAYLOAD_MASS_KG_)
        AS avg_payload_mass_kg FROM
        SPACEXTBL
WHERE   BOOSTER_VERSION = 'F9v1.1'
```

- Result

avg_payload_mass_kg
2928

First Successful Ground Landing Date

- Observed that the dates of the first successful landing outcome on ground pad was 22nd December 2015
- Using the MIN() function to find out the earliest date in the column DATE.
- In the WHERE clause, filter the dataset to perform a search only if Landing outcome is Success (ground pad).

- Query

```
SELECT MIN(DATE)
      AS first_successful_landing_date FROM
      SPACEXTBL
WHERE  LANDING_OUTCOME
      = 'Success (ground pad)'
```

- Result

first_successful_landing_date
2015-12-22

Successful Drone Ship Landing with Payload between 4000 and 6000

- In the WHERE clause, filter the dataset to perform a search if Landing outcome is Success (drone ship)
- Using the AND operator to display a record if additional condition PAYLOAD_MASS_KG_ is between 4000 and 6000.

- Query

```
SELECT BOOSTER_VERSION  
FROM SPACEXTBL  
WHERE LANDING OUTCOME = 'Success (drone ship)' AND  
      (PAYLOAD_MASS_KG_ BETWEEN 4000 AND 6000)
```

- Result

booster_version
F9 FT B1022
F9 FT B1026
F9 FT B1021.2
F9 FT B1031.2

Total Number of Successful and Failure Mission Outcomes

- Using the COUNT() function to calculate the total number of columns.
- Using the GROUP BY statement, groups rows that have the same values into summary rows to find the total number in each Mission_outcome.
- According to the result, SpaceX seems to have successfully completed nearly 99% of its missions.

- Query

```
SELECT MISSION_OUTCOME,  
       COUNT(*) AS total_number  
FROM SPACEXTBL  
GROUP BY MISSION_OUTCOME
```

- Result

mission_outcome	total_number
Failure (in flight)	1
Success	99
Success (payload status unclear)	1

Boosters Carried Maximum Payload

- Using a subquery, first, find the maximum value of the payload by using MAX() function, and second, filter the dataset to perform a search if PAYLOAD_MASS_KG_ is the maximum value of the payload.
- According to the result, version F9 B5 B10xx.x
- boosters could carried the maximum payload.

- Query

```
SELECT DISTINCT BOOSTER_VERSION,  
                PAYLOAD_MASS_KG_  
FROM   SPACEXTBL  
WHERE  PAYLOAD_MASS_KG_ = (  
        SELECT MAX(PAYLOAD_MASS_KG_) FROM  
              SPACEXTBL)
```

- Result

booster_version	payload_mass_kg_
F9 B5 B1048.4	15600
F9 B5 B1048.5	15600
F9 B5 B1049.4	15600
F9 B5 B1049.5	15600
F9 B5 B1049.7	15600
F9 B5 B1051.3	15600
F9 B5 B1051.4	15600
F9 B5 B1051.6	15600
F9 B5 B1056.4	15600
F9 B5 B1058.3	15600
F9 B5 B1060.2	15600
F9 B5 B1060.3	15600

2015 Launch Records

- In the WHERE clause, filter the dataset to perform a search if Landing outcome is Failure (drone ship)
- Using the AND operator to display a record if additional condition YEAR is 2015
- In 2015, there were two landing failures on drone ships.

- Query

```
SELECT  LANDING  OUTCOME,  
        BOOSTER_VERSION,  
        LAUNCH_SITE  
FROM    SPACEXTBL  
WHERE   LANDING OUTCOME  
        = 'Failure (drone ship)' AND  
        YEAR(DATE) = '2015'
```

- Result

landing__outcome	booster_version	launch_site
Failure (drone ship)	F9 v1.1 B1012	CCAFS LC-40
Failure (drone ship)	F9 v1.1 B1015	CCAFS LC-40

Rank Landing Outcomes Between 2010-06-04 and 2017-03-20

- In the WHERE clause, filter the dataset to perform a search if the date is between 2010-06-04 and 2017-03-20.
- Using the ORDER BY keyword to sort the records by total number of landing, and using DESC keyword to sort the records in descending order.
- According to the results, the number of successes and failures between 2010-06-04 and 2017-03-20 was similar.
- Query

```
SELECT LANDING_OUTCOME,  
       COUNT(LANDING_OUTCOME) AS  
       total_number FROM SPACEXTBL  
WHERE  DATE BETWEEN '2010-06-04' AND '2017-  
03-20' GROUP BY LANDING_OUTCOME  
ORDER  BY total_number DESC
```

- Result

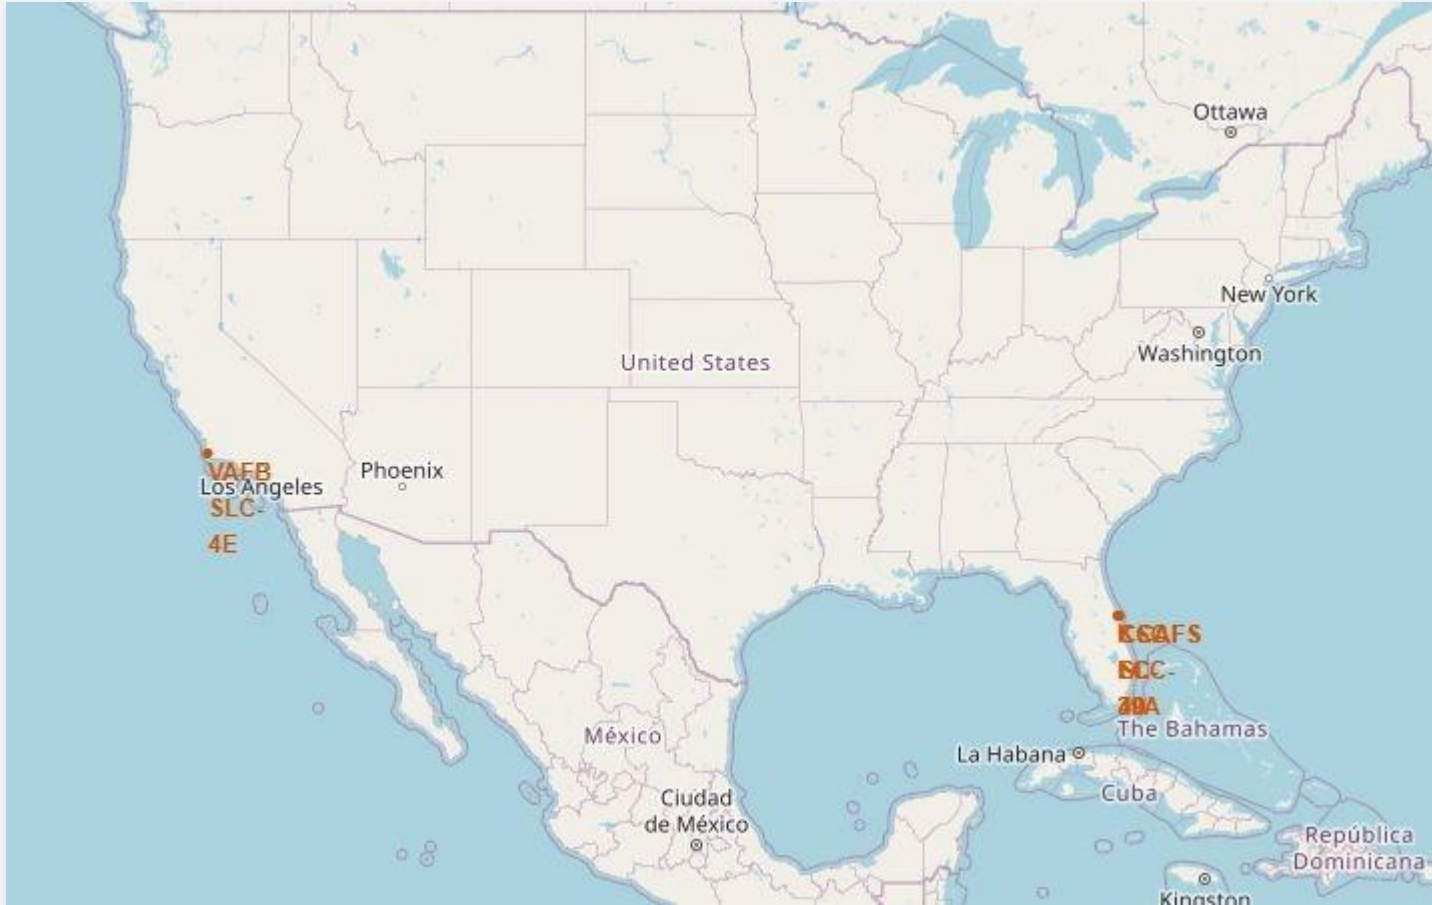
landing__outcome	total_number
No attempt	10
Failure (drone ship)	5
Success (drone ship)	5
Controlled (ocean)	3
Success (ground pad)	3
Failure (parachute)	2
Uncontrolled (ocean)	2
Precluded (drone ship)	1

A satellite view of Earth from space, showing the curvature of the planet and city lights at night. The background is a deep blue gradient.

Section 3

Launch Sites Proximities Analysis

All Launch Sites Locations with map markers







- The left map shows all SpaceX launch sites, and the right map also shows that all launch sites are in the United States coasts
- As can be seen on the map, all launch sites are near the coast.

Markers showing launch sites with color labels



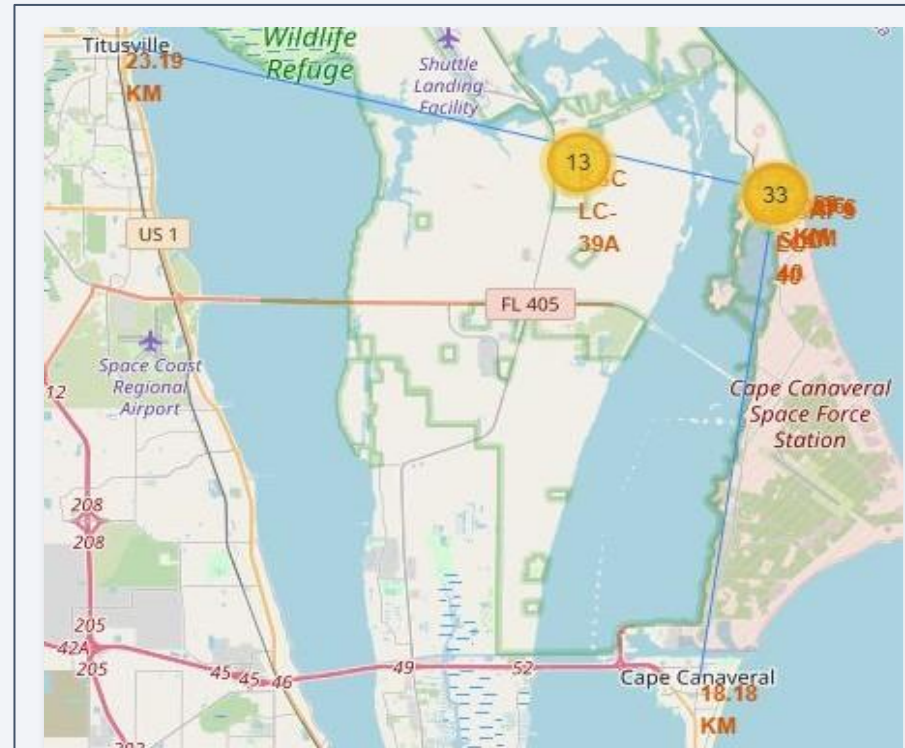
- By clicking on the marker clusters, successful landing (green) or failed landing (red) are displayed.



Proximities of Launch Sites



Are launch sites in close proximity to railways? **Yes**
Are launch sites in close proximity to highways? **Yes**
Are launch sites in close proximity to coastline? **Yes**



Do launch sites keep certain distance away from cities? **Yes**

- It can be found that the launch site is **close to railways and highways** for transportation of equipment or personnel, and is also **close to coastline** and relatively **far from the cities** so that launch failure does not pose a threat.




Section 4

Build a Dashboard with Plotly Dash


Total Success Launches By all sites

- KSLC-39A records the most launch success among all sites.
- The VAFB SLC-4E has the fewest launch success, possibly because
 - the data sample is small, or
 - because it is the only site located in California, so the launch difficulty on the west coast may be higher than on the east coast.



Launch Site with Highest Launch Success Ratio

- KSLC-39A has the highest success rate with 10 landing successes (76.9%) and 3 landing failures (23.1%).



Payload vs. Launch Outcome Scatter Plot for all sites



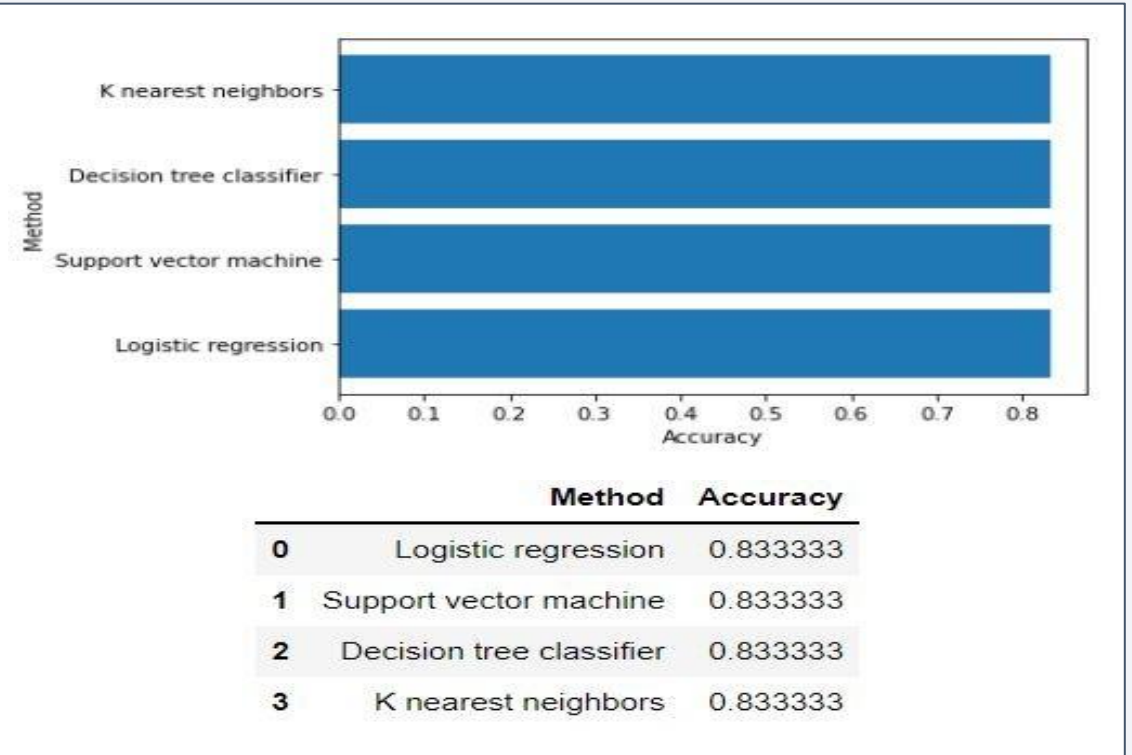
- These figures show that the launch success rate (class 1) for low weighted payloads(0-5000 kg) is higher than that of heavy weighted payloads(5000-10000 kg).

Section 5

Predictive Analysis (Classification)

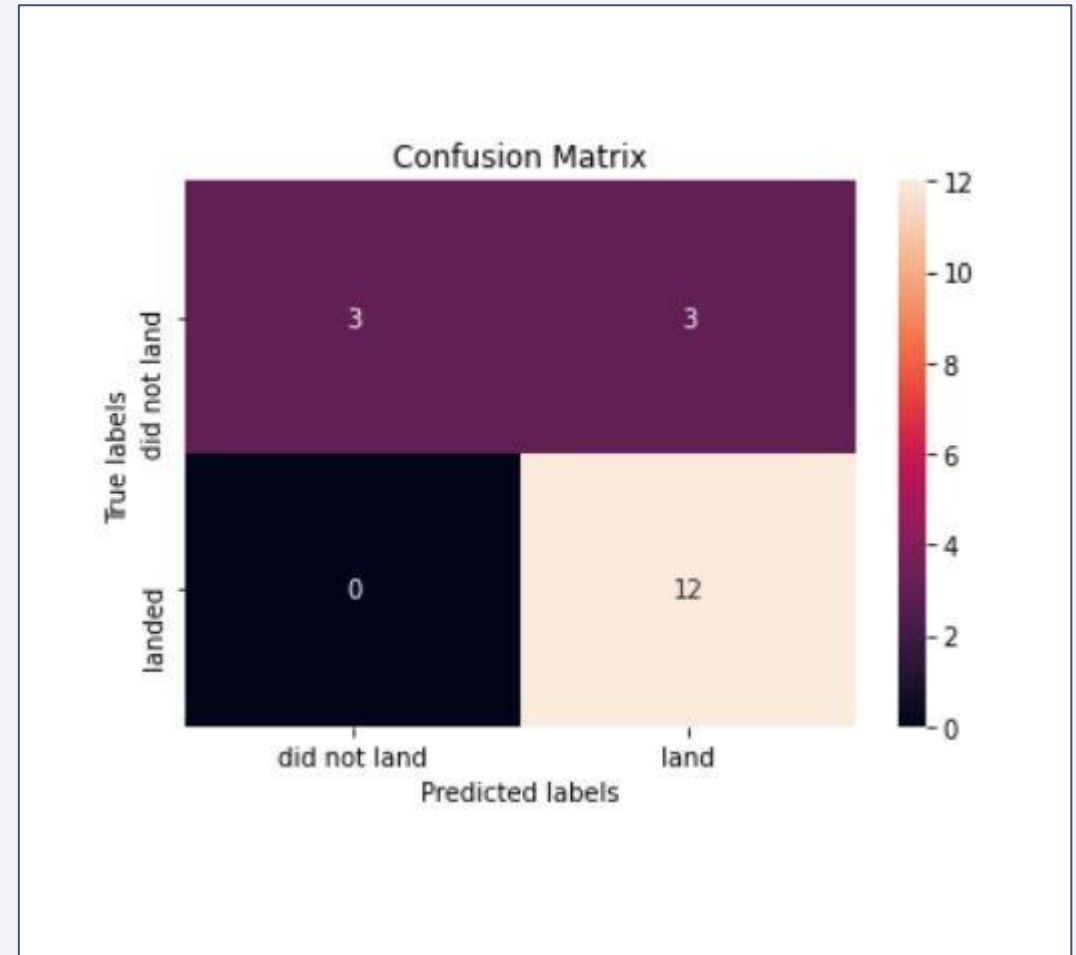
Classification Accuracy

- In the test set, the accuracy of all models was virtually the same at 83.33%.
- It should be noted that the test size was small at 18
- More data is needed to determine the optimal model



Confusion Matrix

- The confusion matrix is the same for all models because all models performed the same for the test set.
- The models predicted 12 successful landings when the true label was successful and 3 failed landings when the true label was failure. But there were also 3 predictions that said successful landings when the true label was failure (false positive).
- Overall, these models predict successful landings.



Conclusions

- As the number of flights increased, the success rate increased, and recently it has exceeded 80%.
- Orbital types SSO, HEO, GEO, and ES-L1 have the highest success rate (100%).
- The launch site is close to railways, highways, and coastline, but far from cities.
- KSLC-39A has the highest number of launch successes and the highest success rate among all sites.
- The launch success rate of low weighted payloads is higher than that of heavy weighted payloads.
- In this dataset, all models have the same accuracy (83.33%), but it seems that more data is needed to determine the optimal model due to the small data size.

Appendix

- <https://github.com/rajesh16rk/IBM-SpaceX-Applied-Capstone-Project>
- [Coursera Applied Data Science Capstone Course URL](#)

Thank you!

